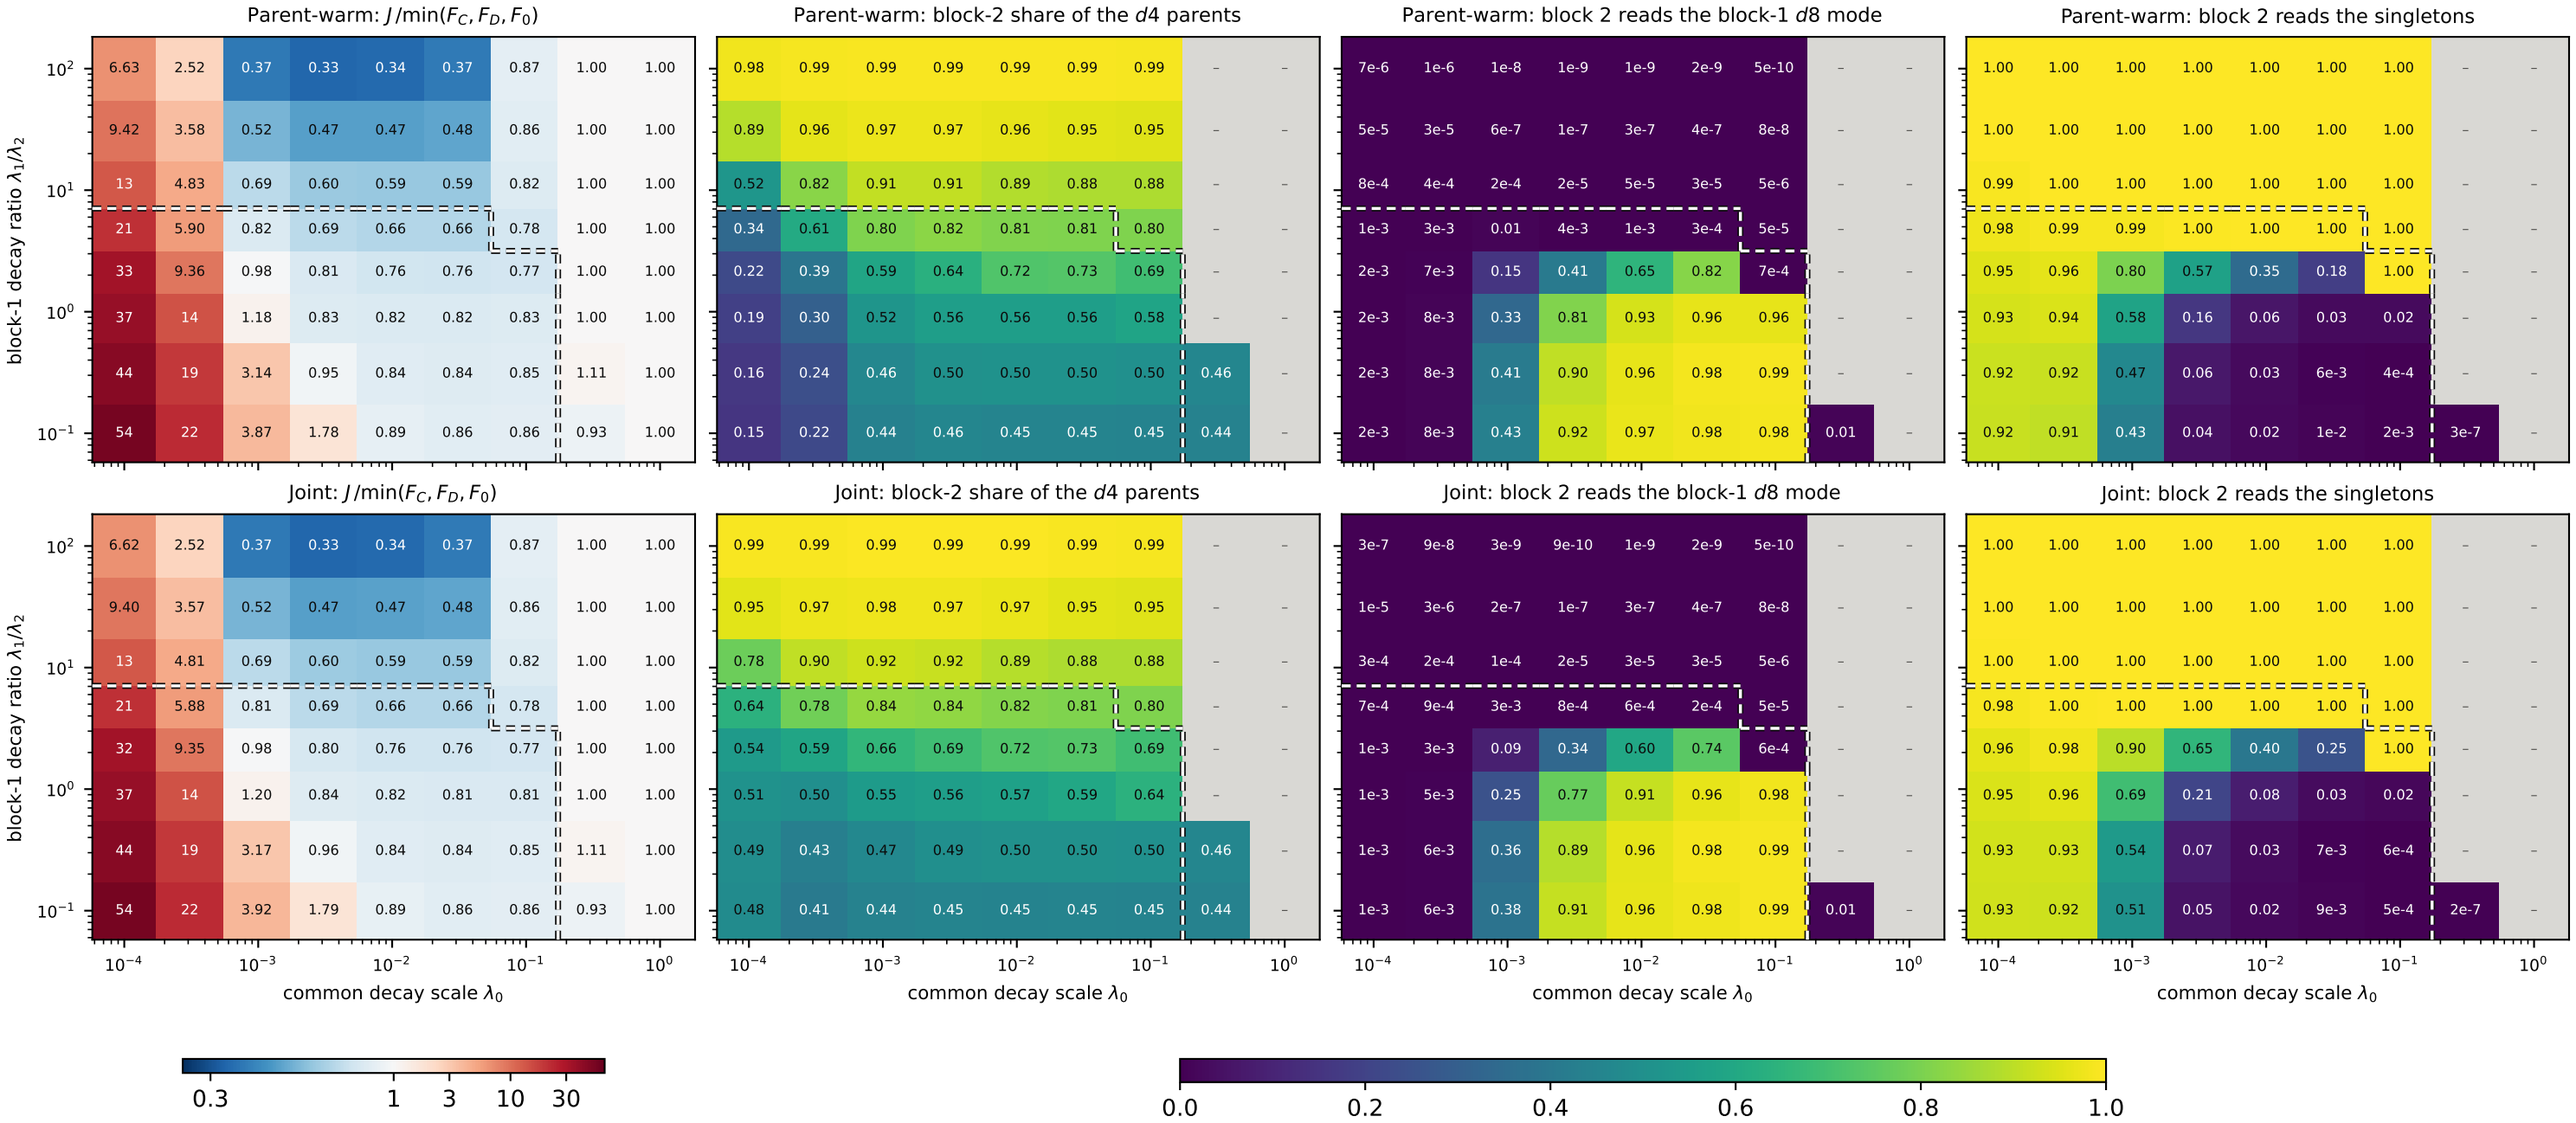Where the parents and the $d8$ output are built: deterministic training (momentum 0.9, 250k steps, hidden lr 0.001, readout init var $1/N$, readout lr 0.001, seed 0)
$J/\min F$: measured objective over the theory's minimal free energy ($F_0 = 3$: zero network); block-2 reads: $|w_{8,n}a_n|$-weighted share of the pre-activation power of block-2 neurons producing $d8$; dashed: theory H_C